

# Information centric networking

# Outline

- ❑ V. Jacobson, D.K. Smetters, J.D. Thornton, M.F. Plass, N.H. Briggs, R.L. Braynard. "Networking named content" , CoNEXT'09.
- ❑ S.K. Fayazbakhsh, Y. Lin, A. Tootoonchian, A. Ghodsiz, T. Koponen, B.M. Maggs, K.C. Ng, V. Sekar, S. Shenker. "Less Pain, Most of the Gain: Incrementally Deployable ICN" , SIGCOMM'13.

# Outline

- ❑ V. Jacobson, D.K. Smetters, J.D. Thornton, M.F. Plass, N.H. Briggs, R.L. Braynard. "Networking named content" , CoNEXT'09.
- ❑ S.K. Fayazbakhsh, Y. Lin, A. Tootoonchian, A. Ghodsiz, T. Koponen, B.M. Maggs, K.C. Ng, V. Sekar, S. Shenker. "Less Pain, Most of the Gain: Incrementally Deployable ICN" , SIGCOMM'13.

# Yesterday's and today's Internet

## 60's and 70's

- ❑ scarce resources; small number of hosts
- ❑ created for resource sharing (e.g., card reader)
- ❑ client/server (point-to-point)
- ❑ static; always connected
- ❑ slow links, slow CPUs, small memory; expensive
- ❑ best effort, no extra features for security, routing

## Now

- ❑ increasing number of devices, some disposable
- ❑ mostly content sharing
- ❑ one produces, many consume; increasing amount of data
- ❑ increasing mobile, potentially intermittent connectivity
- ❑ ever fast links/CPU, abundant storage; cheap
- ❑ increased dependency, increased security threats

**Claim:** Internet communication model based on problems and technologies from 60's, 70's, which today is inadequate

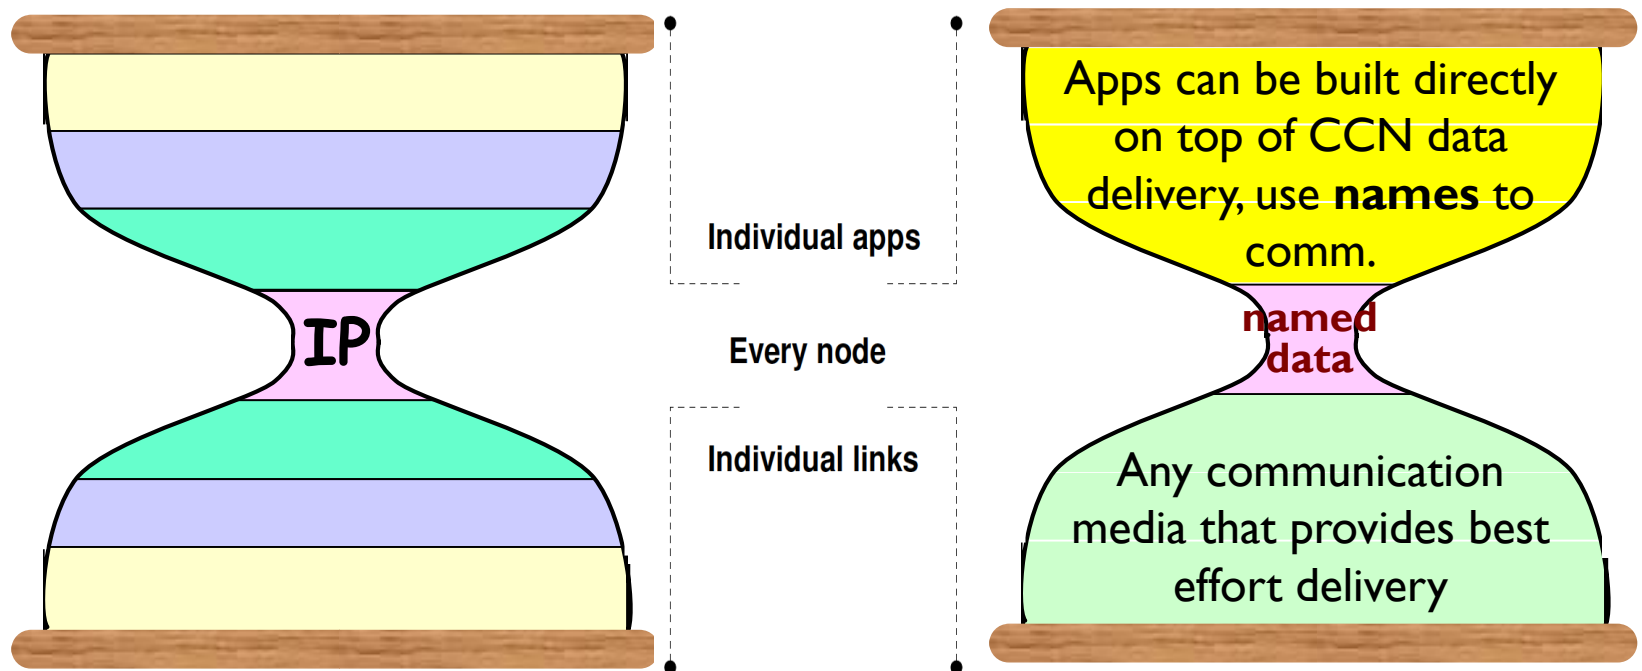
# What needs to be changed?

- ❑ delivery to IP addresses
- ❑ point-to-point delivery
- ❑ lack of built-in security

**How to make change?**

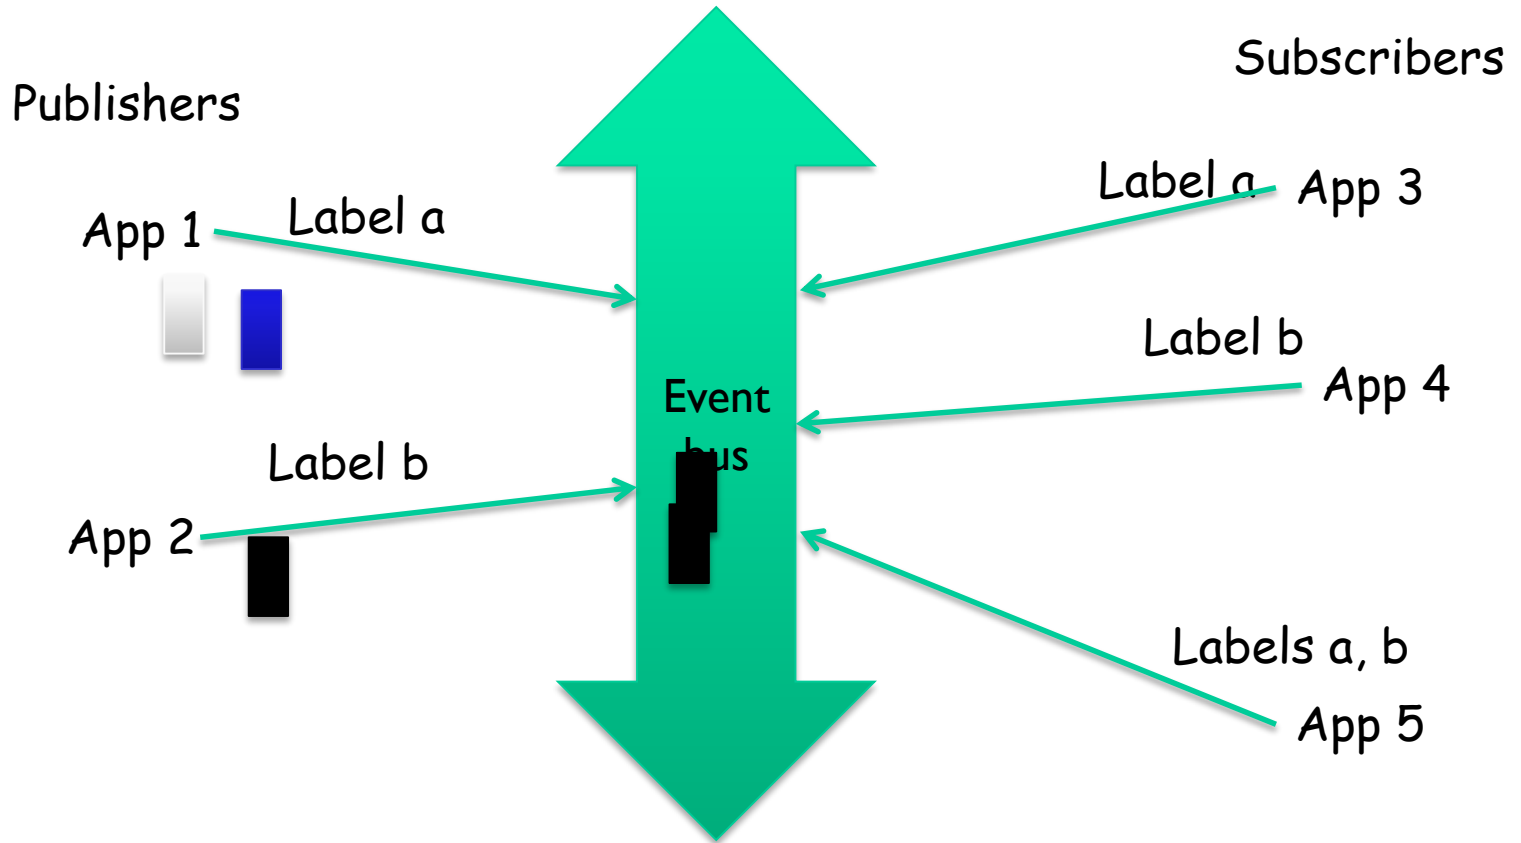
# Content centric networks

Today's Internet delivers packets to destination IP addresses




CCN moves universal component in Internet protocol stack from IP packets to **named data**

# Digression: Publish/Subscribe Model



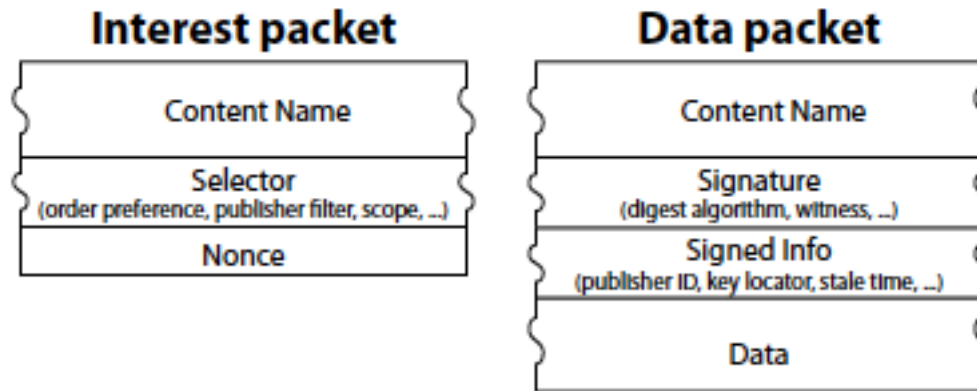
CCN requires content to be requested explicitly  
Pub/Sub model has been proposed as enhancement

# Problem with point to point model





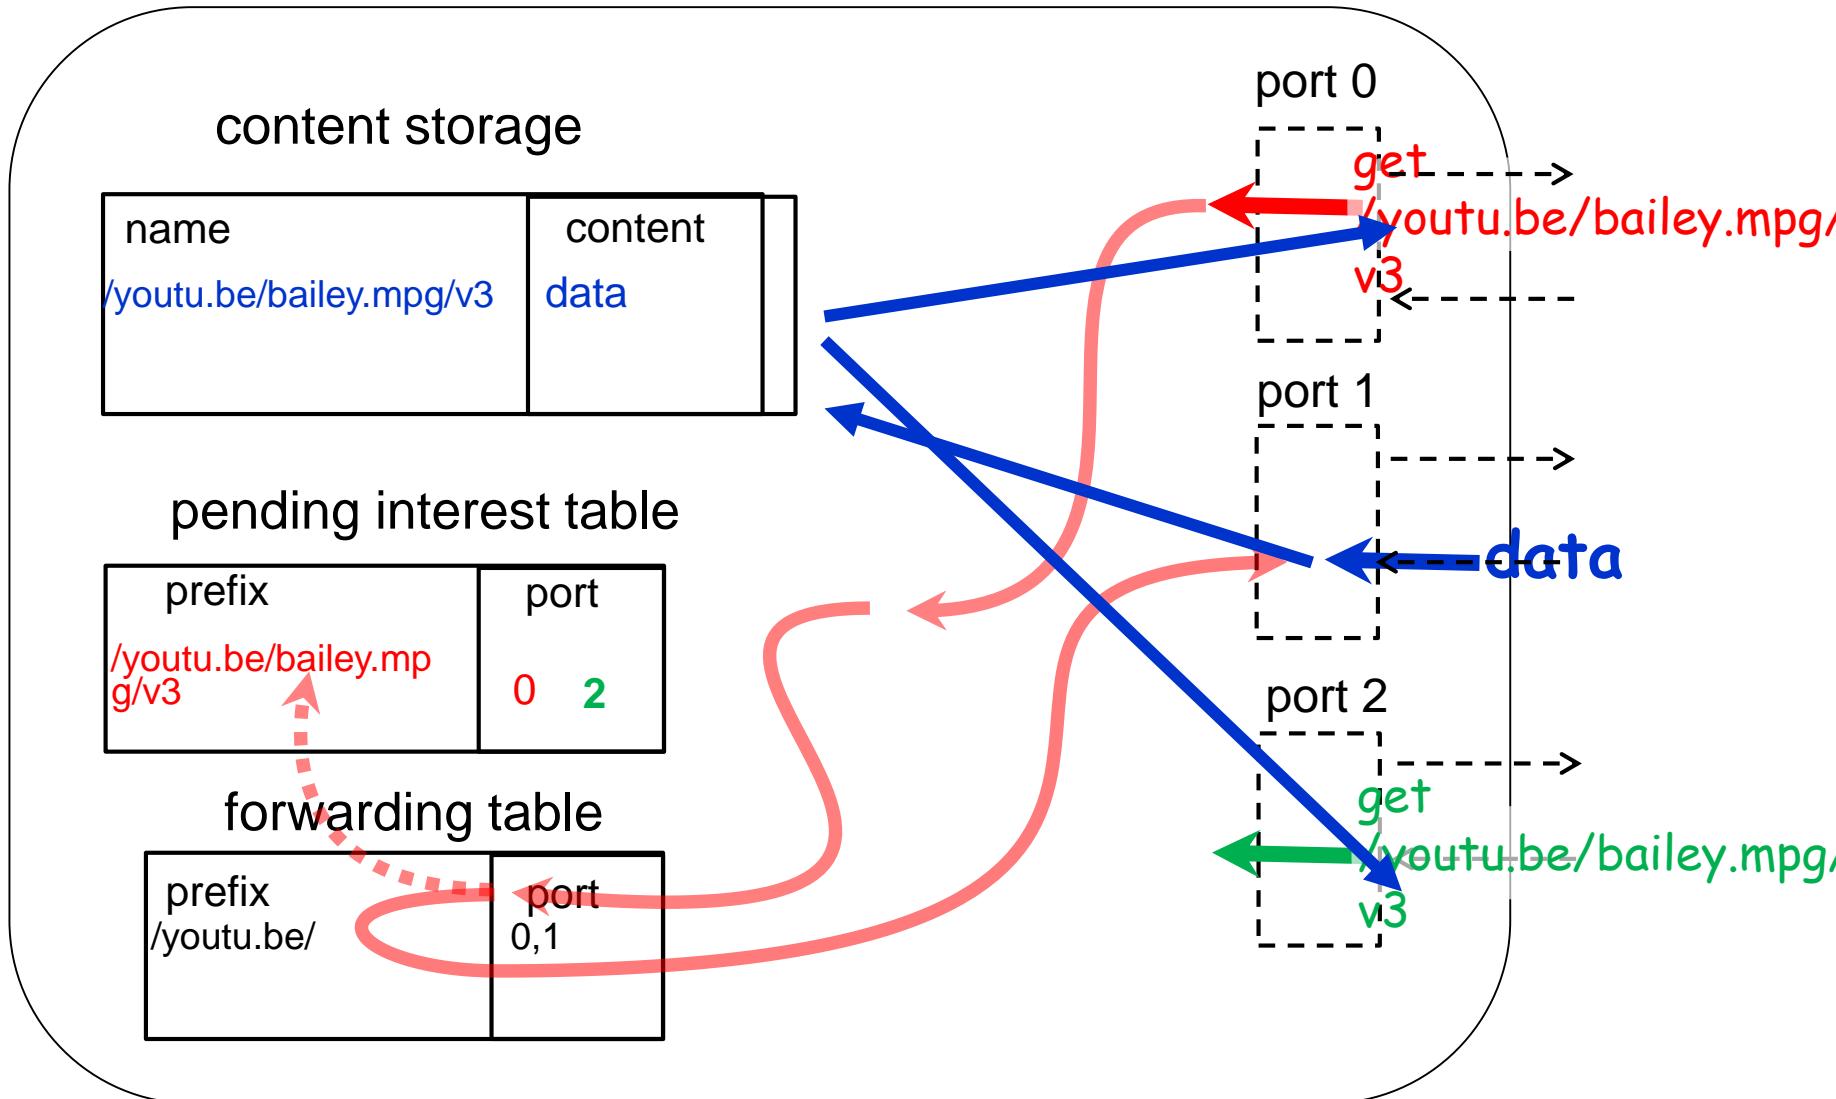
# CCN architecture




## CCN Packets:

- ❑ consumers send **interest packets**, and nodes that can satisfy those interests respond with **data packets**
- ❑ **hierarchical** and **context-dependent** name prefixes (e.g., /local/Friends)
- ❑ **nonces** to prevent Interests from looping

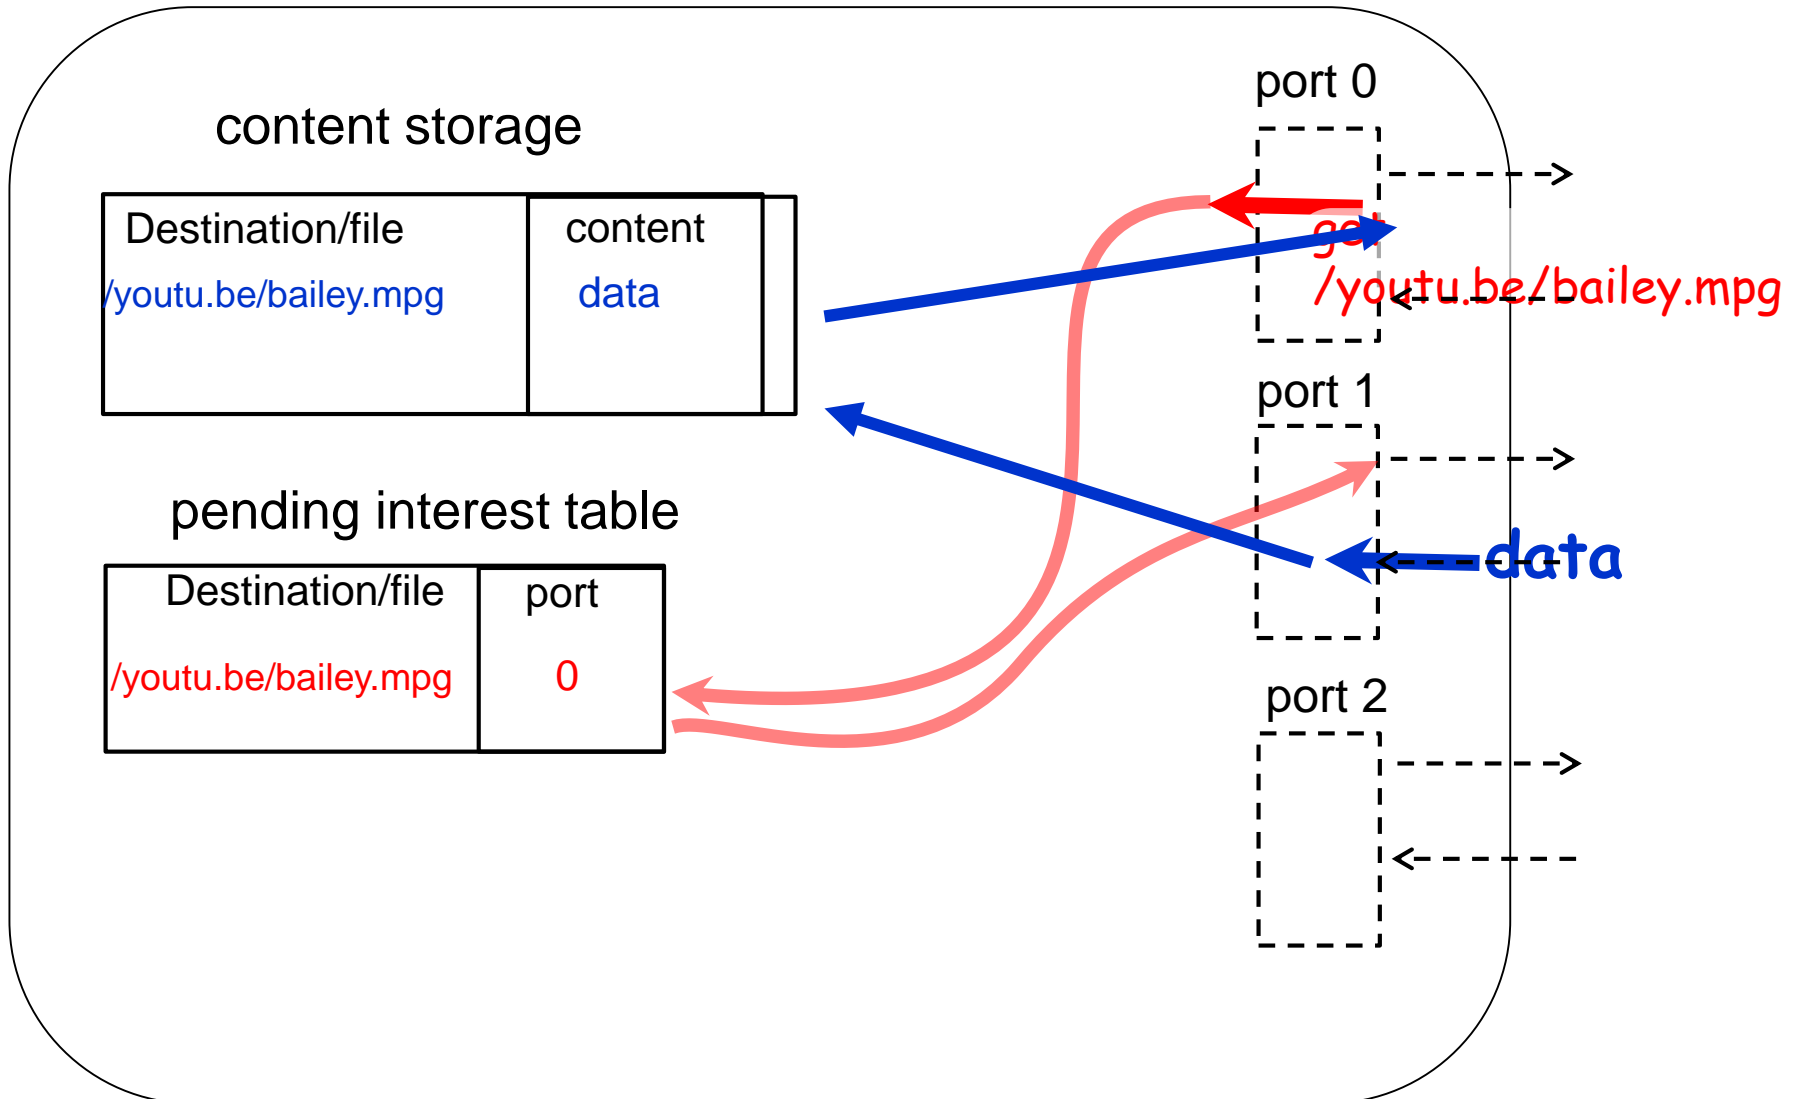
# Routing functionality (Jacobson)




# IP router



# Web Cache (Proxy caching)



# Pending interest table and “bread crumbs”




What happens if both A and B both have content?  
If router sends interests on more than one face, first data pkt returned; some bandwidth wasted

# Some remarks

- ❑ stateful routers
- ❑ interest routed, not data (breadcrumbs)
- ❑ duplicate data packets are discarded
- ❑ **nonces** (random numbers) prevent Interest packets from looping
- ❑ content store uses favorite replacement scheme
- ❑ pending interest entries have timeouts

# CCN enables scalable data dissemination




# Ad hoc networking, mobility, DTN

- ❑ two or more mobile nodes can start communicating with each other as soon as they can physically reach each other


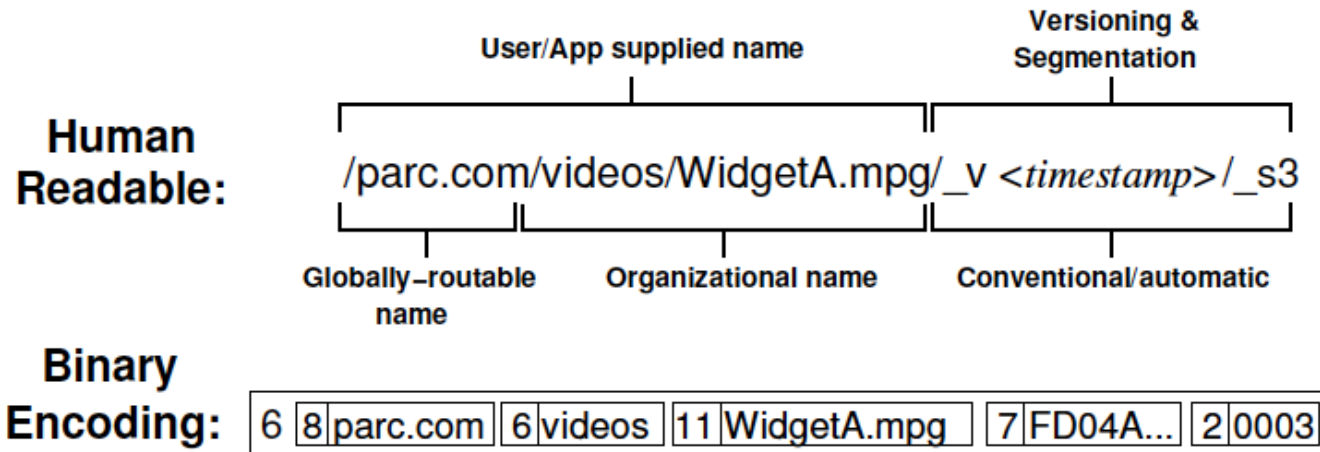


- ❑ CCN provides efficient streaming to mobiles on the move





# Addressing Scheme



- ❑ hierarchical names (components)
- ❑ sequencing

What is maximum height/width of tree?  
How long does look-up take?

# TCP-like features

## Reliability:

- ❑ application resends requests, more flexible  
E.g., app can implement network coding
- ❑ similar to TCP SACK

## Flow Control:

- ❑ at most one data packet per interest packet
- ❑ TCP window advertisements → interest packets

# Other layers

**Strategy layer** (program in Forwarding Table describing how to use faces)

- ❑ multiple interfaces allowed
  - ❖ sendToAll (broadcast), sendToBest (opportunistic routing)

## **Routing:**

- ❑ any routing scheme that works well for IP
  - ❖ IP and CCN forwarding are based on prefixes
- ❑ multi-sources, multi-destinations
- ❑ compatible with IP-based routers (CCN route announcements discarded)

# CCN Security Model


- ❑ IP: point-to-point, secures the channel
- ❑ CCN: secures data, not its container
  - ❖ first, data must be visible in the architecture
  - ❖ then, secure data:
    - associate key with each name, sign data together with name at creation
    - can verify integrity and provenance independent from where it came

# Interest flooding attack

**Flooding:** *Generate large number of interest packets to overwhelm content source*

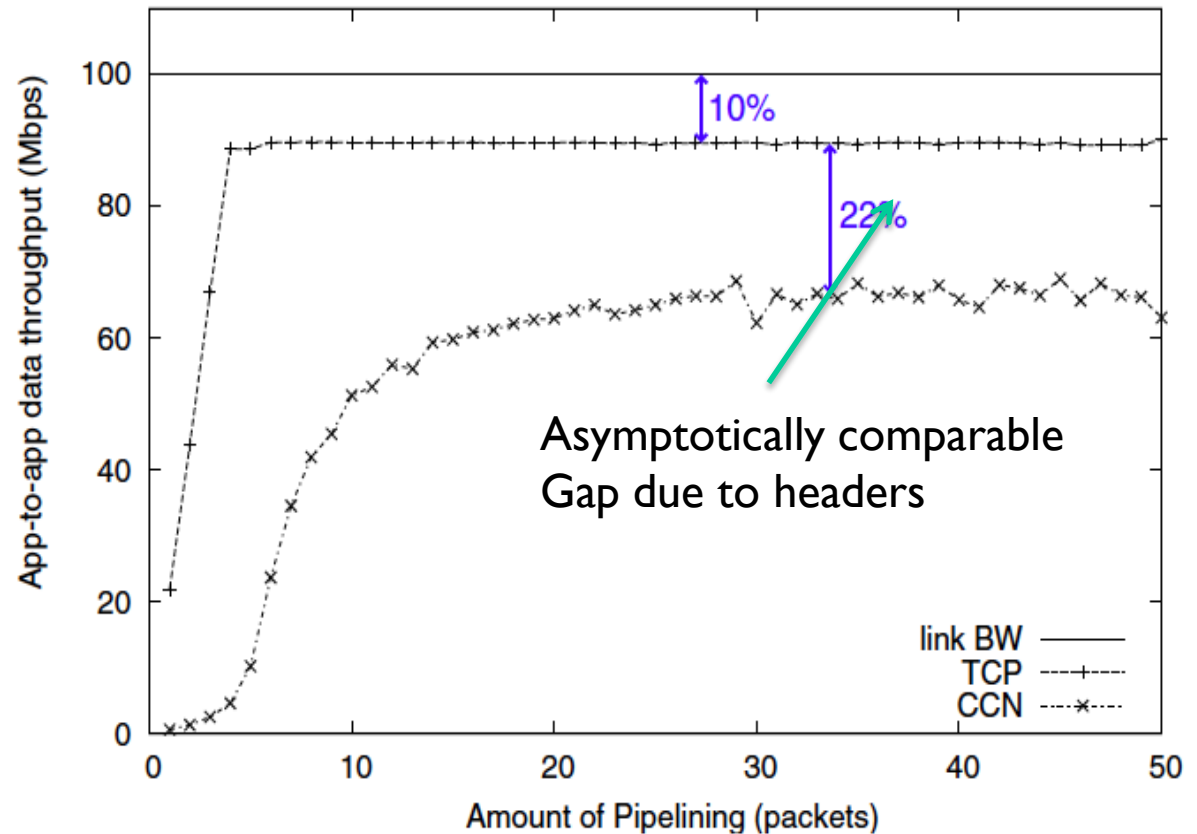

Defenses:

- ❑ nodes can monitor how many interest pkts of same prefix were successfully resolved
- ❑ domain can ask **downstream routers** to throttle number of interests they forward of same prefix

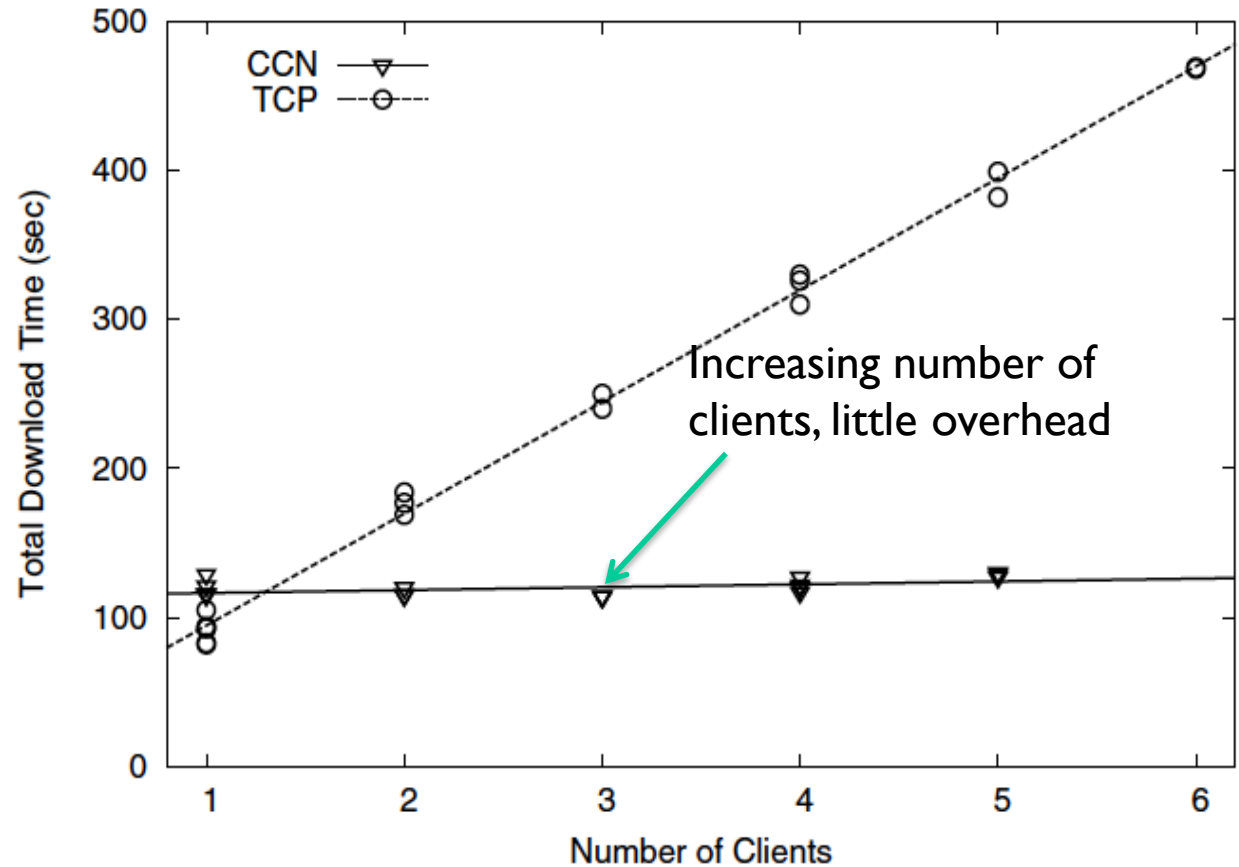



# Experiment I

## □ performance of CCN vs. tcp



# Experiment 2: “multicast” performance



Implicit assumption - Sinks sync'd

# Summary

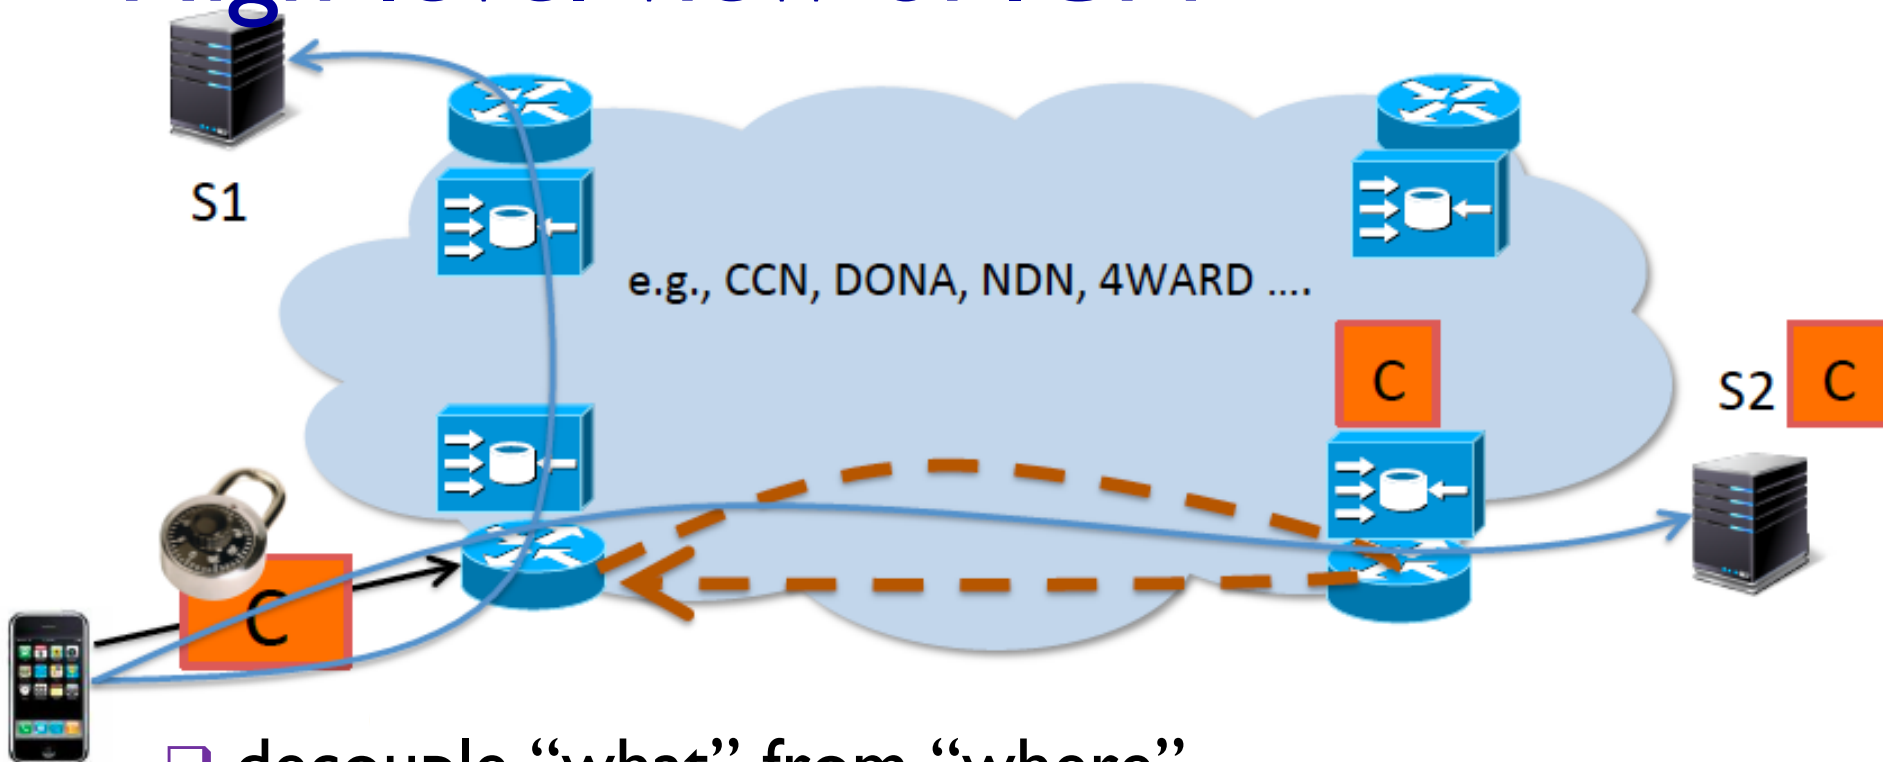
- ❑ CCN - clean-slate architecture for content-based network service
- ❑ based on successes and lessons from today's Internet
- ❑ built-in security, multicast and multipath
- ❑ components to facilitate mobility, ad hoc and disruption-tolerant networks
- ❑ incrementally deployable, but nodes in “bridged” CCN-capable ISPs won't see benefits
- ❑ supports consumer mobility
  - provider mobility?



# Outline

- V. Jacobson, D.K. Smetters, J.D. Thornton, M.F. Plass, N.H. Briggs, R.L. Braynard. "Networking named content" , CoNEXT'09.
- S.K. Fayazbakhsh, Y. Lin, A. Tootoonchian, A. Ghodsiz, T. Koponen, B.M. Maggs, K.C. Ng, V. Sekar, S. Shenker. "Less Pain, Most of the Gain: Incrementally Deployable ICN" , SIGCOMM'13.

# High-level view of ICN



- ❑ decouple “what” from “where”
- ❑ bind content names to content
- ❑ equip network with content caches
- ❑ route based on content names  
e.g.: find nearest replica

# Motivation for work

## Gains

- ☐ lower latency
- ☐ reduced congestion
- ☐ support for mobility
- ☐ intrinsic security

Can we achieve ICN gains without pains?  
e.g., existing technologies?  
e.g., incrementally deployable?

## Pains

- ☐ routers need to be upgraded with caches
- ☐ routing needs to be content based

# Approach: Attribute gains to architectural features

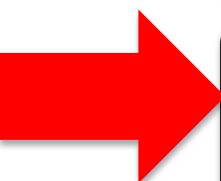
## Quantitative

## Qualitative

- ☐ lower latency
- ☐ reduced congestion

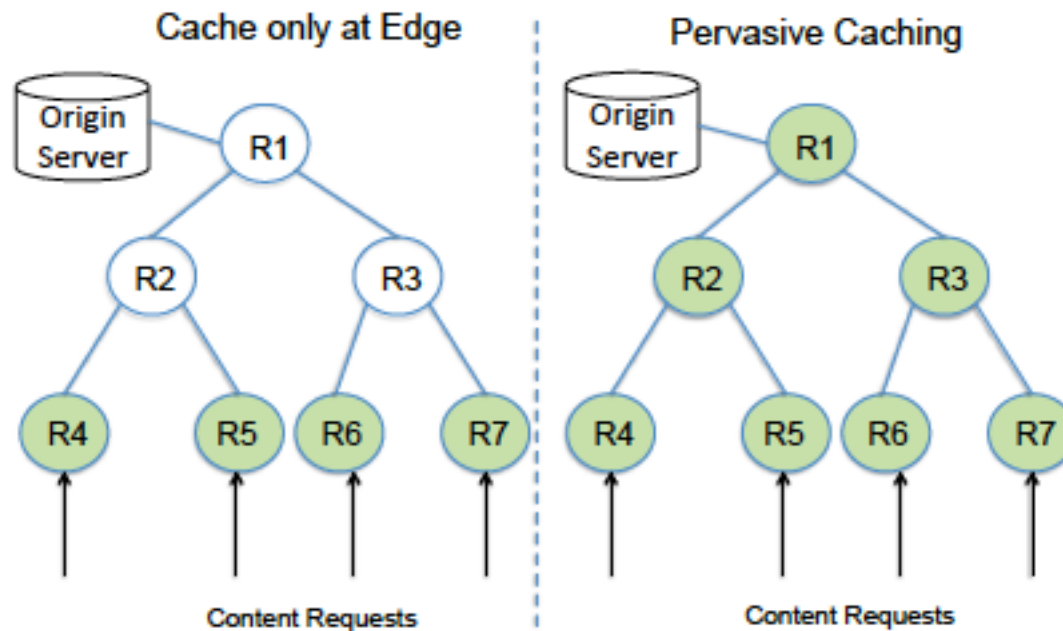
- ☐ support for mobility
- ☐ intrinsic security

- ☐ decouple “what” from “where”
- ☐ bind content names to intent

- 
- ☐ equip network with content caches
  - ☐ route based on content names

# Representative designs

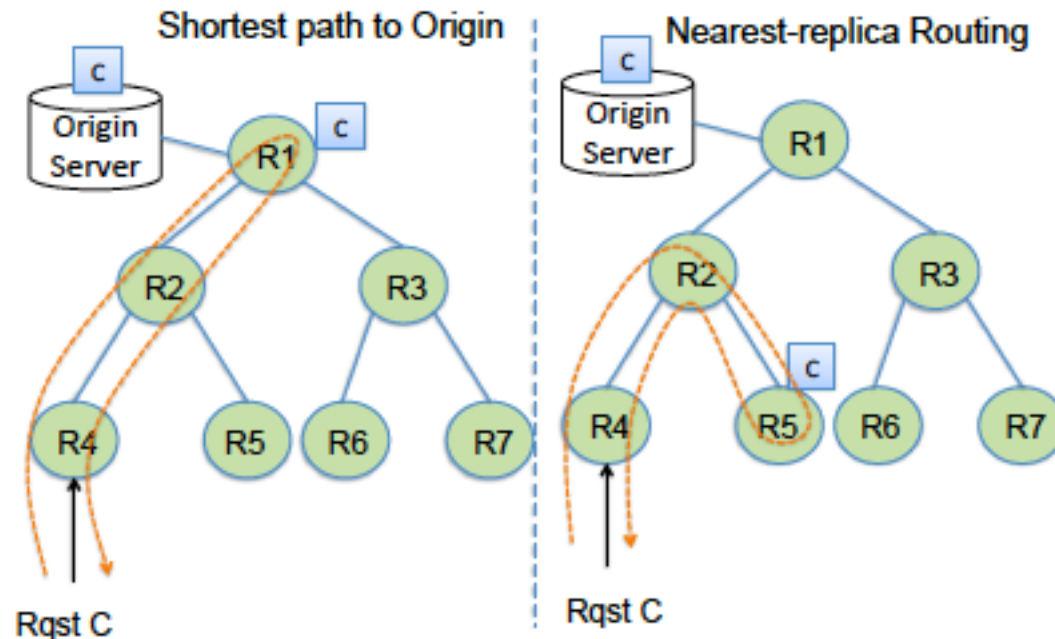
- ❑ *take-away*: **Improvements** on unicast transmissions largely **due to caching**
- ❑ two key dimensions to *design space*:
  - I) cache placement: **Edge** vs. **Pervasive** (everywhere)



# Representative designs

2) How to route requests

Shortest path to origin vs. **Nearest replica**



How is CCN (previous paper) classified?


Pervasive caching and nearest replica routing (?)

# Heavy-tailed workloads

Heavy-tailed: informally, values much larger than average happen significantly often

$X_k$  - popularity of  $k$ -th most popular file.

Zipf's law:  $X_k \sim \frac{1}{k^\alpha}$ ,  $\alpha > 0$



(a) US

# Key takeaways

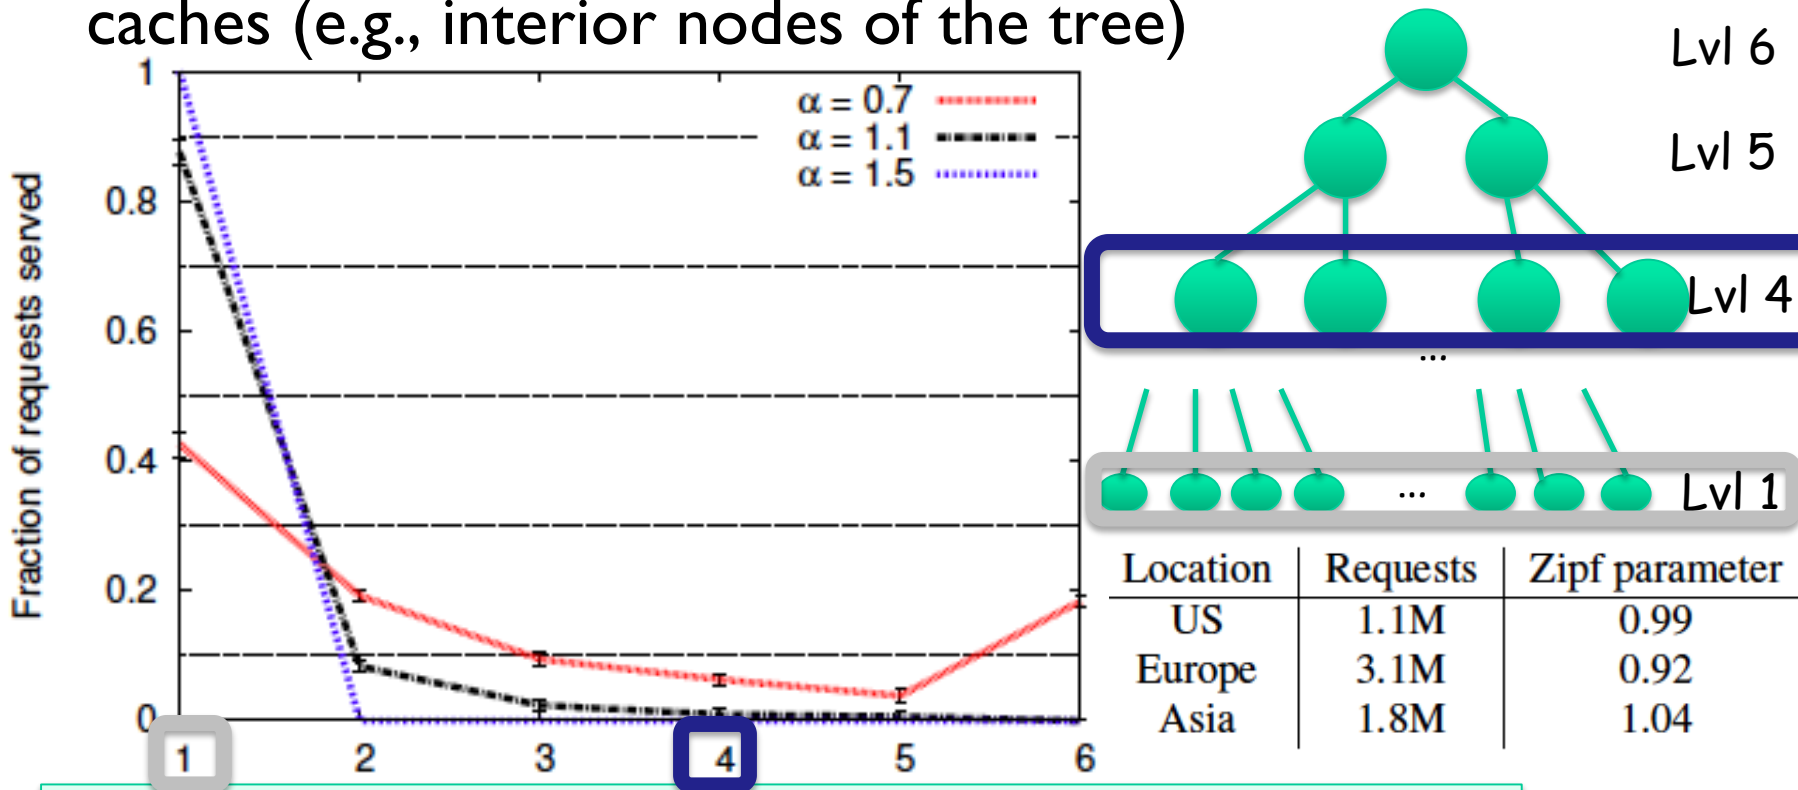
- ❑ to achieve quantitative benefits:
  - cache at “edge”
  - with Zipf-like workloads, pervasive caching and nearest-replica routing don't add much
- ❑ to achieve qualitative benefits:
  - build on HTTP

Basis for incrementally deployable ICN



# Heavy-tailed workloads: implications

- ❑ caching a few of most popular items yields large hit ratios
- ❑ larger exponent  $\alpha$ , faster popularity decays
- ❑ decreasing improvement from setting extra nodes as caches (e.g., interior nodes of the tree)



Take-away: caching at edges suffices

# Simulation setup

Edge

Real CDN  
request logs

Cache provisioning  
~ 5% of objects  
Uniform requests


LRU replacement

Access tree

PoP-level topologies (Rocketfuel) augmented with access trees

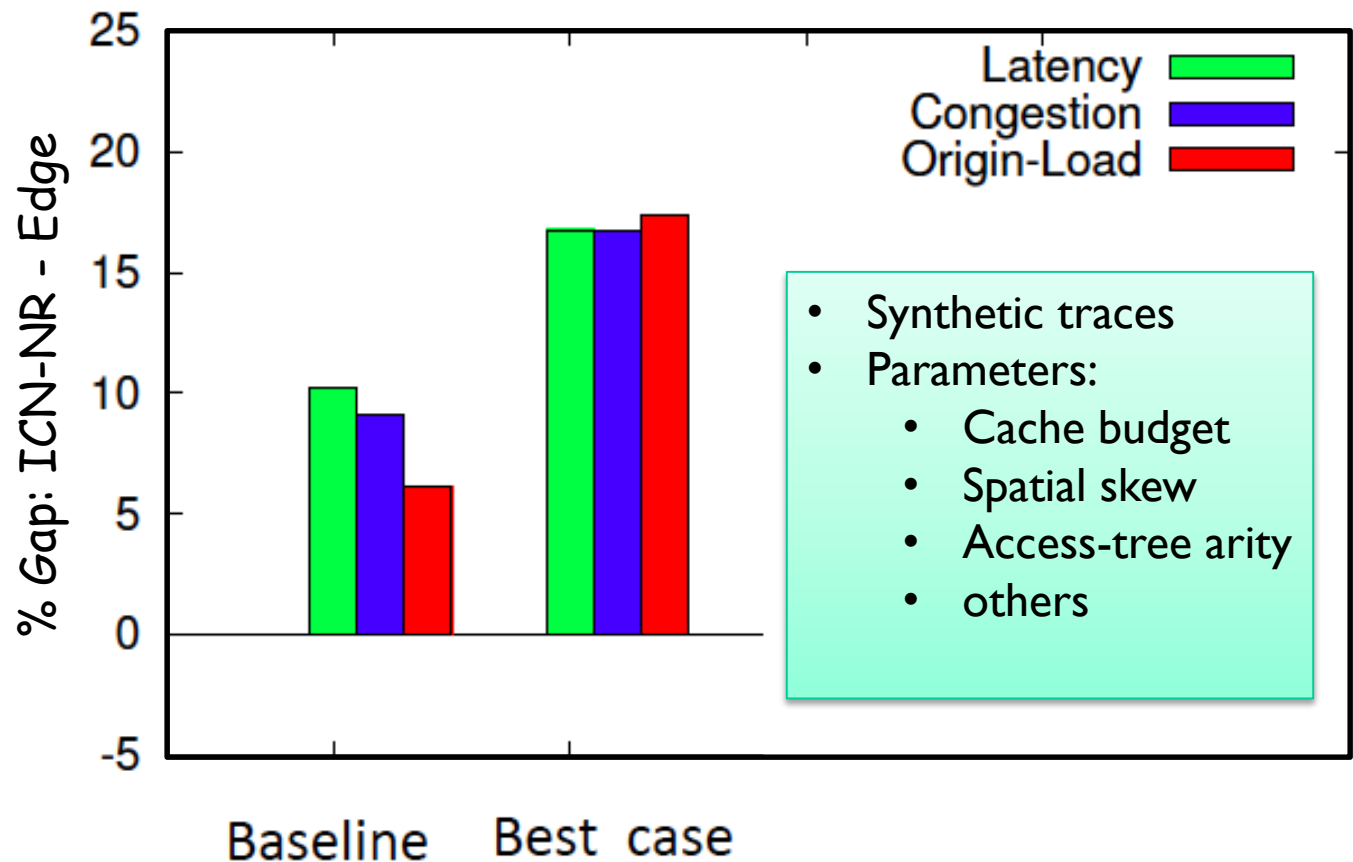
Assume name-based routing, lookup incurs zero cost

# Request latency



Gap between architectures small ( $< 10\%$ )  
Similar results for congestion + server load

# Sensitivity Analysis



Little difference; in best case, ICN-NR only 17% better  
Gap can be “easily” reduced

E.g., normalized budget or cooperative strategies

# Implications of Edge Caching

- ❑ incrementally deployable
  - ❖ domains get benefits without relying on others
- ❑ incentive deployable
  - ❖ domains' users get benefits if domain deploys caches

# Revisiting Qualitative Aspects

## 1. Decouple names from locations

### Build on HTTP

- Can be viewed as providing “get-by-name” abstraction
- Can reuse existing web protocols (e.g., proxy discovery)

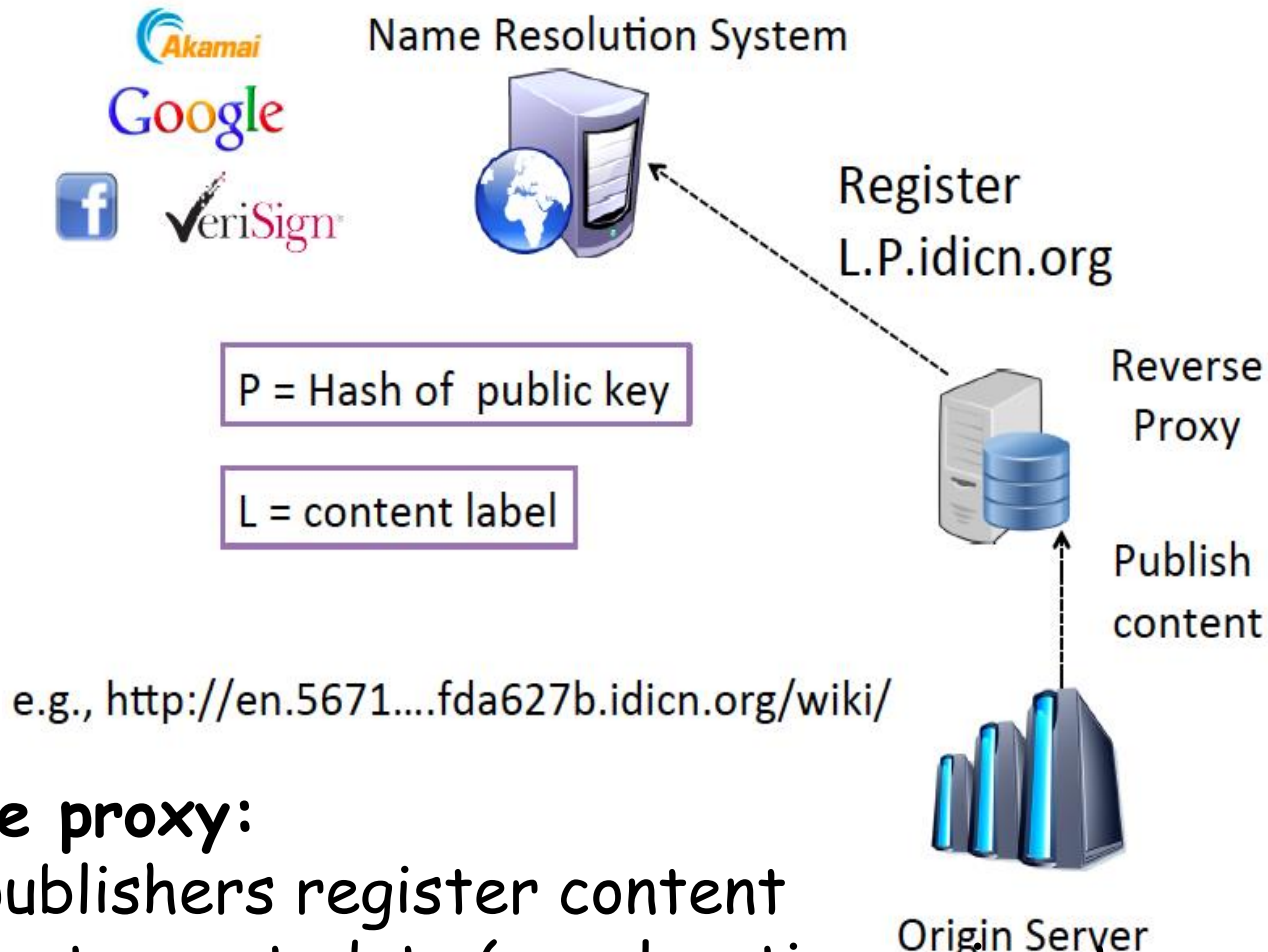
## 2. Binding names to intents

### Use self-certifying names

e.g., “Magnet” URI schemes

Extend HTTP for “crypto” and other metadata

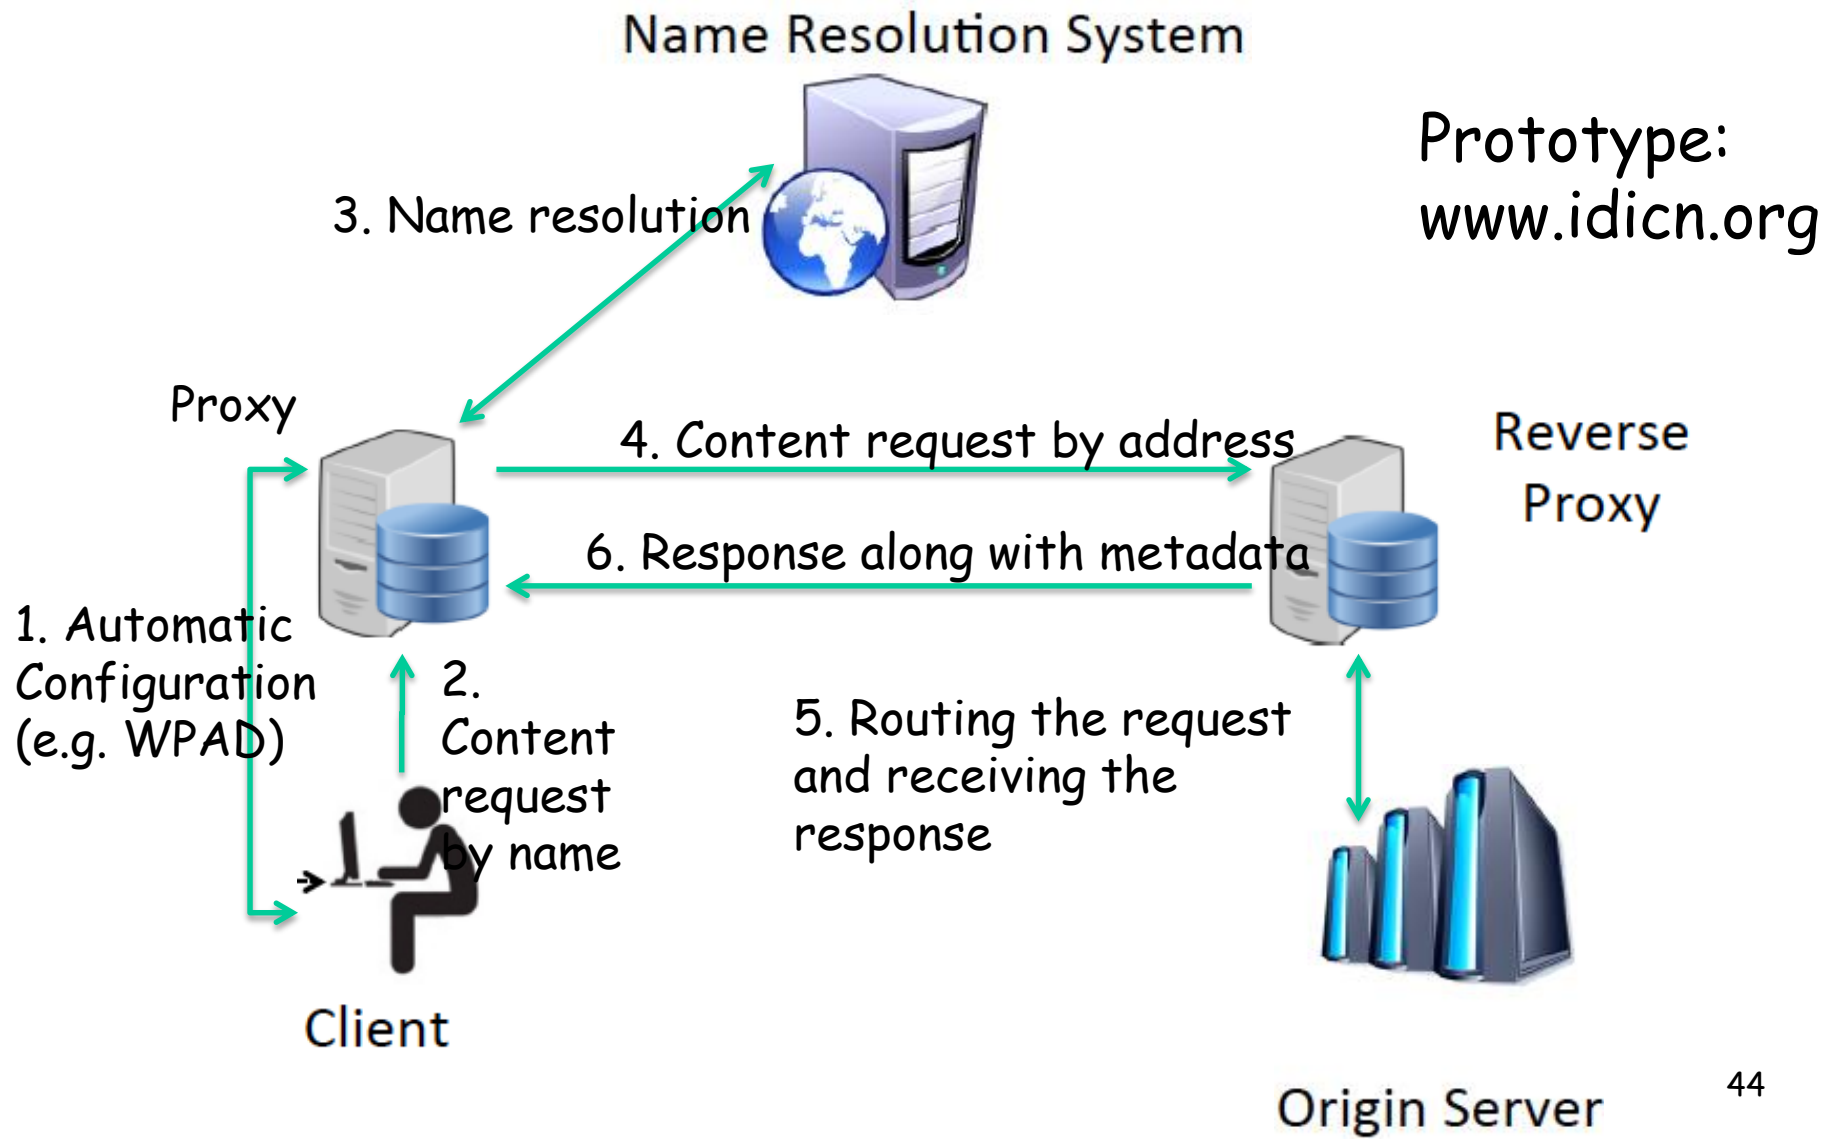
# idICN: Content Registration



## Reverse proxy:

- Let publishers register content
- Generates metadata (e.g., locations, signatures, policies)
- Receives requests by name and return content + metadata

# idICN: Content Delivery





# Summary

- ❑ gains of ICN with less pain
  - ❖ latency, congestion, security
  - ❖ without changes to routers or routing
- ❑ quantitative benefits with “edge” solutions
  - ❖ **pervasive caching, nearest-replica** routing **not needed**
- ❑ qualitative benefits with existing techniques
  - ❖ existing HTTP + HTTP-based extensions
  - ❖ incrementally deployable
- ❑ idlCN: one possible feasible realization
  - ❖ open issues: economics, other benefits, future workloads
- ❑ no multicast, support for mobility

# Quick comparison

## CCN

- ❖ clean-slate
- ❖ requires changing routers
- ❖ pervasive caching, nearest replica routing
- ❖ multiple source-destinations
- ❖ built-in security; protection against DoS attacks

## idICN

- ❖ based on existing infrastructure/protocols
- ❖ edge caching, cooperative routing requests
- ❖ point-to-point
- ❖ security thru extending HTTP to negotiate metadata and standardizing self-certifying naming scheme

**Both cases, produce networks of caches.  
Evaluation involves understanding interaction  
between caches**